

Random_forests

December 10, 2018

1 Random Forests

```
In [38]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

from IPython.display import set_matplotlib_formats
set_matplotlib_formats('svg')

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, balanced_accuracy_score, make_scorer
from sklearn.metrics import roc_auc_score, roc_curve
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.model_selection import cross_val_score

In [39]: # Get features
samples = pd.read_csv("samples.csv", index_col=0)
microbes = pd.read_csv("microbes.csv", index_col=0)
metabolites = pd.read_csv("metabolites.csv", index_col=0)
combined_features = pd.concat([microbes, metabolites], axis=1)
# Label vector
labels = samples.case
```

1.1 Metabolites

```
In [64]: # We'll use 80% for training, and 20% for testing
X_train, X_test, labels_train, labels_test = train_test_split(
    metabolites, labels, test_size=0.2, random_state=42)
```

Random forest

```
In [65]: rf = RandomForestClassifier(n_estimators=10, random_state=42)
#scores = cross_val_score(rf, X_train, y_train, cv=5)
#print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))
scores = cross_val_score(rf, X_train, labels_train, cv=5,
                        scoring=make_scorer(roc_auc_score))
print("AUC: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))
```

AUC: 0.63 (+/- 0.20)

```
In [66]: # Train the best model with the entire training set
rf.fit(X_train, labels_train)
microbes_pred = rf.predict(X_test)
balanced_accuracy = balanced_accuracy_score(labels_test, microbes_pred)
accuracy = accuracy_score(labels_test, microbes_pred)
auc = roc_auc_score(y_true=labels_test, y_score=[j for i, j in rf.predict_proba(X_test)
print("Accuracy score:", accuracy)
print("Balanced accuracy score:", balanced_accuracy)
print("AUC score:", auc)
```

Accuracy_score: 0.851851851852

Balanced accuracy score: 0.75

AUC score: 0.799342105263

1.2 Generate hyperparameter grid

```
In [67]: # Number of trees in random forest
n_estimators = [int(x) for x in np.linspace(start = 10, stop = 2000, num = 10)]
# Number of features to consider at every split
max_features = ['auto', 'sqrt']
# Maximum number of levels in tree
max_depth = [int(x) for x in np.linspace(2, 110, num = 11)]
max_depth.append(None)
# Minimum number of samples required to split a node
min_samples_split = [2, 5, 10]
# Minimum number of samples required at each leaf node
min_samples_leaf = [1, 2, 4]
# Method of selecting samples for training each tree
bootstrap = [True, False]
```

```
# Create the random grid
random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf,
               'bootstrap': bootstrap,
               'random_state': [42]}
```

random_grid

```
Out[67]: {'n_estimators': [10, 231, 452, 673, 894, 1115, 1336, 1557, 1778, 2000],
          'max_features': ['auto', 'sqrt'],
          'max_depth': [2, 12, 23, 34, 45, 56, 66, 77, 88, 99, 110, None],
          'min_samples_split': [2, 5, 10],
```

```

'min_samples_leaf': [1, 2, 4],
'bootstrap': [True, False],
'random_state': [42]}

```

```

In [68]: # Use the random grid to search for best hyperparameters
# First create the base model to tune
rf = RandomForestClassifier()
# Random search of parameters, using 3 fold cross validation,
# search across 100 different combinations, and use all available cores
rf_random = RandomizedSearchCV(estimator = rf, param_distributions = random_grid,
                               n_iter = 100, cv = 5, verbose=2, random_state=42,
                               n_jobs = -1, scoring=make_scorer(roc_auc_score))

# Fit the random search model
rf_random.fit(X_train, labels_train)

```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 50.6s
[Parallel(n_jobs=-1)]: Done 154 tasks    | elapsed: 3.8min
[Parallel(n_jobs=-1)]: Done 357 tasks    | elapsed: 7.6min
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 10.3min finished
/home/ravila/Software/miniconda/lib/python3.6/site-packages/sklearn/model_selection/_search.py
DeprecationWarning)

```

```

Out [68]: RandomizedSearchCV(cv=5, error_score='raise-deprecating',
                             estimator=RandomForestClassifier(bootstrap=True, class_weight=None, criteri
                             max_depth=None, max_features='auto', max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators='warn', n_jobs=None,
                             oob_score=False, random_state=None, verbose=0,
                             warm_start=False),
                             fit_params=None, iid='warn', n_iter=100, n_jobs=-1,
                             param_distributions={'n_estimators': [10, 231, 452, 673, 894, 1115, 1336, 1
                             pre_dispatch='2*n_jobs', random_state=42, refit=True,
                             return_train_score='warn', scoring=make_scorer(roc_auc_score),
                             verbose=2)

```

```

In [69]: rf_random.best_params_

```

```

Out [69]: {'random_state': 42,
           'n_estimators': 10,
           'min_samples_split': 2,
           'min_samples_leaf': 1,
           'max_features': 'sqrt',

```

```
'max_depth': 45,  
'bootstrap': True}
```

```
In [70]: rf_random.best_score_
```

```
Out[70]: 0.62751831501831501
```

Default model

```
In [71]: base_model = RandomForestClassifier(n_estimators = 10, random_state = 42)  
base_model.fit(X_train, labels_train)  
pred = base_model.predict(X_test)
```

```
balanced_accuracy = balanced_accuracy_score(labels_test, pred)  
accuracy = accuracy_score(labels_test, pred)  
auc = roc_auc_score(y_true=labels_test, y_score=[j for i, j in base_model.predict_proba(X_test)])  
print("Accuracy score:", accuracy)  
print("Balanced accuracy score:", balanced_accuracy)  
print("AUC score:", auc)
```

```
Accuracy_score: 0.851851851852
```

```
Balanced accuracy score: 0.75
```

```
AUC score: 0.799342105263
```

```
In [72]: rf_best = rf_random.best_estimator_  
rf_best.fit(X_train, labels_train)  
pred = rf_best.predict(X_test)
```

```
balanced_accuracy = balanced_accuracy_score(labels_test, pred)  
accuracy = accuracy_score(labels_test, pred)  
auc = roc_auc_score(y_true=labels_test, y_score=[j for i, j in rf_best.predict_proba(X_test)])  
print("Accuracy score:", accuracy)  
print("Balanced accuracy score:", balanced_accuracy)  
print("AUC score:", auc)
```

```
Accuracy_score: 0.851851851852
```

```
Balanced accuracy score: 0.75
```

```
AUC score: 0.799342105263
```

```
In [73]: # Plot an ROC curve
```

```
fpr, tpr, _ = roc_curve(y_true=labels_test, y_score=[j for i, j in rf_best.predict_proba(X_test)])  
auc = roc_auc_score(y_true=labels_test, y_score=[j for i, j in rf_best.predict_proba(X_test)])  
plt.plot(fpr, tpr, label="AUC = {0:.2f}".format(auc))  
plt.plot([0, 1], [0, 1], color='red', lw=1, linestyle='--')  
plt.xlim([0.0, 1.0])  
plt.ylim([0.0, 1.0])  
plt.xlabel('False Positive Rate')
```

```
plt.ylabel('True Positive Rate')  
plt.legend()  
plt.show()
```

