

NeuralNet

December 10, 2018

1 Neural Networks

```
In [3]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

from IPython.display import set_matplotlib_formats
set_matplotlib_formats('svg')

from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, balanced_accuracy_score, make_scorer
from sklearn.metrics import roc_auc_score, roc_curve
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.model_selection import cross_val_score

In [21]: # Get features
samples = pd.read_csv("samples.csv", index_col=0)
microbes = pd.read_csv("microbes.csv", index_col=0)
metabolites = pd.read_csv("metabolites.csv", index_col=0)
combined_features = pd.concat([microbes, metabolites], axis=1)
# Label vector
labels = samples.case
```

1.1 Metabolites

```
In [28]: # We'll use 80% for training, and 20% for testing
X_train, X_test, labels_train, labels_test = train_test_split(
    metabolites, labels, test_size=0.2, random_state=42)

In [97]: nnet = MLPClassifier(solver='lbfgs', max_iter=300, learning_rate='invscaling',
                             hidden_layer_sizes=(10), activation='relu', random_state=42)
#scores = cross_val_score(rf, X_train, y_train, cv=5)
#print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))
scores = cross_val_score(nnet, X_train, labels_train, cv=5,
                          scoring=make_scorer(roc_auc_score))
print("AUC: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))

AUC: 0.76 (+/- 0.12)
```

```
In [156]: mnet = MLPClassifier(solver='lbfgs', alpha=.0001, max_iter=200, learning_rate='invscale',
                               hidden_layer_sizes=(10), activation='relu', random_state=42)
#scores = cross_val_score(rf, X_train, y_train, cv=5)
#print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))
scores = cross_val_score(mnet, X_train, labels_train, cv=5,
                          scoring=make_scorer(roc_auc_score))
print("AUC: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))
```

AUC: 0.76 (+/- 0.12)

1.2 Generate hyperparameter grid

```
In [157]: hidden_layer_sizes = [(10, ), (20, )]
activation = ['identity', 'logistic', 'tanh', 'relu']
solver = ['lbfgs', 'sgd', 'adam']
alpha = [0.0001, .001, .01, .1]
batch_size = ['auto']
learning_rate = ['constant', 'invscaling', 'adaptive']
learning_rate_init = [0.001]
max_iter = [200, 300]
validation_fraction = [.2]
epsilon = [1e-08]
n_iter_no_change = [10]

# Create the random grid
random_grid = {'hidden_layer_sizes': hidden_layer_sizes,
               'activation': activation,
               'solver': solver,
               'alpha': alpha,
               'validation_fraction': validation_fraction,
               'learning_rate': learning_rate,
               'learning_rate_init': learning_rate_init,
               'max_iter': max_iter,
               'n_iter_no_change': n_iter_no_change,
               'random_state': [42]}

random_grid
```

```
Out[157]: {'hidden_layer_sizes': [(10, ), (20, )],
          'activation': ['identity', 'logistic', 'tanh', 'relu'],
          'solver': ['lbfgs', 'sgd', 'adam'],
          'alpha': [0.0001, 0.001, 0.01, 0.1],
          'validation_fraction': [0.2],
          'learning_rate': ['constant', 'invscaling', 'adaptive'],
          'learning_rate_init': [0.001],
          'max_iter': [200, 300],
```

```
'n_iter_no_change': [10],
'random_state': [42]}
```

```
In [170]: # Use the random grid to search for best hyperparameters
# First create the base model to tune
mlp = MLPClassifier()
# Random search of parameters, using 3 fold cross validation,
# search across 100 different combinations, and use all available cores
mlp_random = RandomizedSearchCV(estimator = mlp, param_distributions = random_grid,
                                n_iter = 100, cv = 5, verbose=2, random_state=42,
                                n_jobs = -1, scoring=make_scorer(roc_auc_score))

# Fit the random search model
mlp_random.fit(X_train, labels_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed:    7.5s
[Parallel(n_jobs=-1)]: Done 154 tasks     | elapsed:   47.6s
[Parallel(n_jobs=-1)]: Done 357 tasks     | elapsed:   1.5min
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed:   2.0min finished
```

```
Out[170]: RandomizedSearchCV(cv=5, error_score='raise-deprecating',
                             estimator=MLPClassifier(activation='relu', alpha=0.0001, batch_size='auto',
                             beta_2=0.999, early_stopping=False, epsilon=1e-08,
                             hidden_layer_sizes=(100,), learning_rate='constant',
                             learning_rate_init=0.001, max_iter=200, momentum=0.9,
                             n_iter_no_change=10, nesterovs_momentum=True, power_t=0.5,
                             random_state=None, shuffle=True, solver='adam', tol=0.0001,
                             validation_fraction=0.1, verbose=False, warm_start=False),
                             fit_params=None, iid='warn', n_iter=100, n_jobs=-1,
                             param_distributions={'hidden_layer_sizes': [(10,), (20,)], 'activation': [
                             pre_dispatch='2*n_jobs', random_state=42, refit=True,
                             return_train_score='warn', scoring=make_scorer(roc_auc_score),
                             verbose=2)
```

```
In [171]: mlp_random.best_params_
```

```
Out[171]: {'validation_fraction': 0.2,
'solver': 'lbfgs',
'random_state': 42,
'n_iter_no_change': 10,
'max_iter': 300,
'learning_rate_init': 0.001,
'learning_rate': 'adaptive',
'hidden_layer_sizes': (10,),
```

```
'alpha': 0.0001,  
'activation': 'relu'}
```

```
In [172]: mlp_random.best_score_
```

```
Out[172]: 0.7601877289377289
```

Default model

```
In [173]: base_model = MLPClassifier(random_state = 42)  
base_model.fit(X_train, labels_train)  
pred = base_model.predict(X_test)
```

```
balanced_accuracy = balanced_accuracy_score(labels_test, pred)  
accuracy = accuracy_score(labels_test, pred)  
auc = roc_auc_score(y_true=labels_test, y_score=[j for i, j in base_model.predict_proba(X_test)])  
print("Accuracy score:", accuracy)  
print("Balanced accuracy score:", balanced_accuracy)  
print("AUC score:", auc)
```

```
Accuracy_score: 0.777777777778  
Balanced accuracy score: 0.697368421053  
AUC score: 0.684210526316
```

```
In [174]: mlp_best = mlp_random.best_estimator_  
mlp_best.fit(X_train, labels_train)  
pred = mlp_best.predict(X_test)
```

```
balanced_accuracy = balanced_accuracy_score(labels_test, pred)  
accuracy = accuracy_score(labels_test, pred)  
auc = roc_auc_score(y_true=labels_test, y_score=[j for i, j in mlp_best.predict_proba(X_test)])  
print("Accuracy score:", accuracy)  
print("Balanced accuracy score:", balanced_accuracy)  
print("AUC score:", auc)
```

```
Accuracy_score: 0.777777777778  
Balanced accuracy score: 0.697368421053  
AUC score: 0.75
```

```
In [175]: # Plot an ROC curve
```

```
fpr, tpr, _ = roc_curve(y_true=labels_test, y_score=[j for i, j in mlp_best.predict_proba(X_test)])  
auc = roc_auc_score(y_true=labels_test, y_score=[j for i, j in mlp_best.predict_proba(X_test)])  
plt.plot(fpr, tpr, label="AUC = {:.2f}".format(auc))  
plt.plot([0, 1], [0, 1], color='red', lw=1, linestyle='--')  
plt.xlim([0.0, 1.0])  
plt.ylim([0.0, 1.0])  
plt.xlabel('False Positive Rate')
```

```
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```

