

# Tune\_Random\_forests

December 10, 2018

## 1 Random Forests

```
In [54]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

from IPython.display import set_matplotlib_formats
set_matplotlib_formats('svg')

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, balanced_accuracy_score, make_scorer
from sklearn.metrics import roc_auc_score, roc_curve
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.model_selection import cross_val_score

In [2]: # Get features
samples = pd.read_csv("samples.csv", index_col=0)
microbes = pd.read_csv("microbes.csv", index_col=0)
metabolites = pd.read_csv("metabolites.csv", index_col=0)
combined_features = pd.concat([microbes, metabolites], axis=1)
# Label vector
labels = samples.case
```

### 1.1 Metabolites

```
In [3]: # We'll use 80% for training, and 20% for testing
X_train, X_test, labels_train, labels_test = train_test_split(
    metabolites, labels, test_size=0.2, random_state=42)
```

Random forest

```
In [37]: rf = RandomForestClassifier(n_estimators=30, min_samples_split=10, min_samples_leaf=1,
                                   max_features='sqrt', max_depth=80, bootstrap=True, random_state=42)

#scores = cross_val_score(rf, X_train, y_train, cv=5)
#print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))
scores = cross_val_score(rf, X_train, labels_train, cv=5,
                        scoring=make_scorer(roc_auc_score))
print("AUC: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))
```

AUC: 0.65 (+/- 0.18)

```
In [39]: rf = RandomForestClassifier(n_estimators=11, min_samples_split=2, min_samples_leaf=1,
                                   max_features='sqrt', max_depth=10, bootstrap=True, random_state=42)
#scores = cross_val_score(rf, X_train, y_train, cv=5)
#print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))
scores = cross_val_score(rf, X_train, labels_train, cv=5,
                        scoring=make_scorer(roc_auc_score))
print("AUC: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))
```

AUC: 0.64 (+/- 0.15)

```
In [ ]: {'random_state': 42,
        'n_estimators': 30,
        'min_samples_split': 10,
        'min_samples_leaf': 1,
        'max_features': 'sqrt',
        'max_depth': 80,
        'bootstrap': True}
```

## 1.2 Generate hyperparameter grid

```
In [40]: # Number of trees in random forest
n_estimators = [int(x) for x in np.linspace(start = 11, stop = 50, num = 30)]
# Number of features to consider at every split
max_features = ['sqrt']
# Maximum number of levels in tree
max_depth = [int(x) for x in np.linspace(2, 80, num = 11)]
max_depth.append(None)
# Minimum number of samples required to split a node
min_samples_split = [2, 5, 10]
# Minimum number of samples required at each leaf node
min_samples_leaf = [1, 2, 4]
# Method of selecting samples for training each tree
bootstrap = [True, False]

# Create the random grid
random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf,
               'bootstrap': bootstrap,
               'random_state': [42]}

random_grid
```

```

Out[40]: {'n_estimators': [11,
    12,
    13,
    15,
    16,
    17,
    19,
    20,
    21,
    23,
    24,
    25,
    27,
    28,
    29,
    31,
    32,
    33,
    35,
    36,
    37,
    39,
    40,
    41,
    43,
    44,
    45,
    47,
    48,
    50],
    'max_features': ['sqrt'],
    'max_depth': [2, 9, 17, 25, 33, 41, 48, 56, 64, 72, 80, None],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False],
    'random_state': [42]}

```

```

In [45]: # Use the random grid to search for best hyperparameters
# First create the base model to tune
rf = RandomForestClassifier()
# Random search of parameters, using 3 fold cross validation,
# search across 100 different combinations, and use all available cores
rf_random = RandomizedSearchCV(estimator = rf, param_distributions = random_grid,
    n_iter = 400, cv = 5, verbose=2, random_state=42,
    n_jobs = -1, scoring=make_scorer(roc_auc_score))

# Fit the random search model
rf_random.fit(X_train, labels_train)

```

Fitting 5 folds for each of 400 candidates, totalling 2000 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 1.7s
[Parallel(n_jobs=-1)]: Done 154 tasks     | elapsed: 8.1s
[Parallel(n_jobs=-1)]: Done 357 tasks     | elapsed: 19.3s
[Parallel(n_jobs=-1)]: Done 640 tasks     | elapsed: 33.8s
[Parallel(n_jobs=-1)]: Done 1005 tasks    | elapsed: 53.9s
[Parallel(n_jobs=-1)]: Done 1450 tasks    | elapsed: 1.3min
[Parallel(n_jobs=-1)]: Done 1977 tasks    | elapsed: 1.8min
[Parallel(n_jobs=-1)]: Done 2000 out of 2000 | elapsed: 1.8min finished
/home/ravila/Software/miniconda/lib/python3.6/site-packages/sklearn/model_selection/_search.py
DeprecationWarning)
```

```
Out [45]: RandomizedSearchCV(cv=5, error_score='raise-deprecating',
                             estimator=RandomForestClassifier(bootstrap=True, class_weight=None, criteri
                             max_depth=None, max_features='auto', max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators='warn', n_jobs=None,
                             oob_score=False, random_state=None, verbose=0,
                             warm_start=False),
                             fit_params=None, iid='warn', n_iter=400, n_jobs=-1,
                             param_distributions={'n_estimators': [11, 12, 13, 15, 16, 17, 19, 20, 21, 2
                             pre_dispatch='2*n_jobs', random_state=42, refit=True,
                             return_train_score='warn', scoring=make_scorer(roc_auc_score),
                             verbose=2)
```

```
In [46]: rf_random.best_params_
```

```
Out [46]: {'random_state': 42,
           'n_estimators': 12,
           'min_samples_split': 5,
           'min_samples_leaf': 1,
           'max_features': 'sqrt',
           'max_depth': 56,
           'bootstrap': True}
```

```
In [47]: rf_random.best_score_
```

```
Out [47]: 0.64938186813186816
```

Default model

```
In [50]: base_model = RandomForestClassifier(n_estimators = 10, random_state = 42)
         base_model.fit(X_train, labels_train)
         pred = base_model.predict(X_test)
```

```

balanced_accuracy = balanced_accuracy_score(labels_test, pred)
accuracy = accuracy_score(labels_test, pred)
auc = roc_auc_score(y_true=labels_test, y_score=[j for i, j in base_model.predict_proba(X_test)])
print("Accuracy_score:", accuracy)
print("Balanced accuracy score:", balanced_accuracy)
print("AUC score:", auc)

```

```

Accuracy_score: 0.851851851852
Balanced accuracy score: 0.75
AUC score: 0.799342105263

```

```

In [51]: rf_best = rf_random.best_estimator_
         rf_best.fit(X_train, labels_train)
         pred = rf_best.predict(X_test)

         balanced_accuracy = balanced_accuracy_score(labels_test, pred)
         accuracy = accuracy_score(labels_test, pred)
         auc = roc_auc_score(y_true=labels_test, y_score=[j for i, j in rf_best.predict_proba(X_test)])
         print("Accuracy_score:", accuracy)
         print("Balanced accuracy score:", balanced_accuracy)
         print("AUC score:", auc)

```

```

Accuracy_score: 0.777777777778
Balanced accuracy score: 0.625
AUC score: 0.815789473684

```

```

In [59]: # Plot an ROC curve
         fpr, tpr, _ = roc_curve(y_true=labels_test, y_score=[j for i, j in rf_best.predict_proba(X_test)])
         auc = roc_auc_score(y_true=labels_test, y_score=[j for i, j in rf_best.predict_proba(X_test)])
         plt.plot(fpr, tpr, label="AUC = {0:.2f}".format(auc))
         plt.plot([0, 1], [0, 1], color='red', lw=1, linestyle='--')
         plt.xlim([0.0, 1.0])
         plt.ylim([0.0, 1.05])
         plt.xlabel('False Positive Rate')
         plt.ylabel('True Positive Rate')
         plt.legend()
         plt.show()

```

